

# Carrera de Ingeniería de Sistemas y Computación

## Programación Paralela

Profiling Code: Detecting Performance Bottlenecks

Mayo 2022

Tek



# Profiling Code

---

- Using Python's built-in profiling module: The `cProfile` module (Python's recommended module to gather profiling information)

`python -m cProfile -s cumulative  
MyPythonCode.py > profile.out`

*execute the `cProfile` module*

*profile statistics ordered by cumulative time*

*Pass the `MyPythonCode.py` script file to the `cProfile` module*

*Redirect the output to the `profile.out` file*

# Profiling Code (cProfile sample)

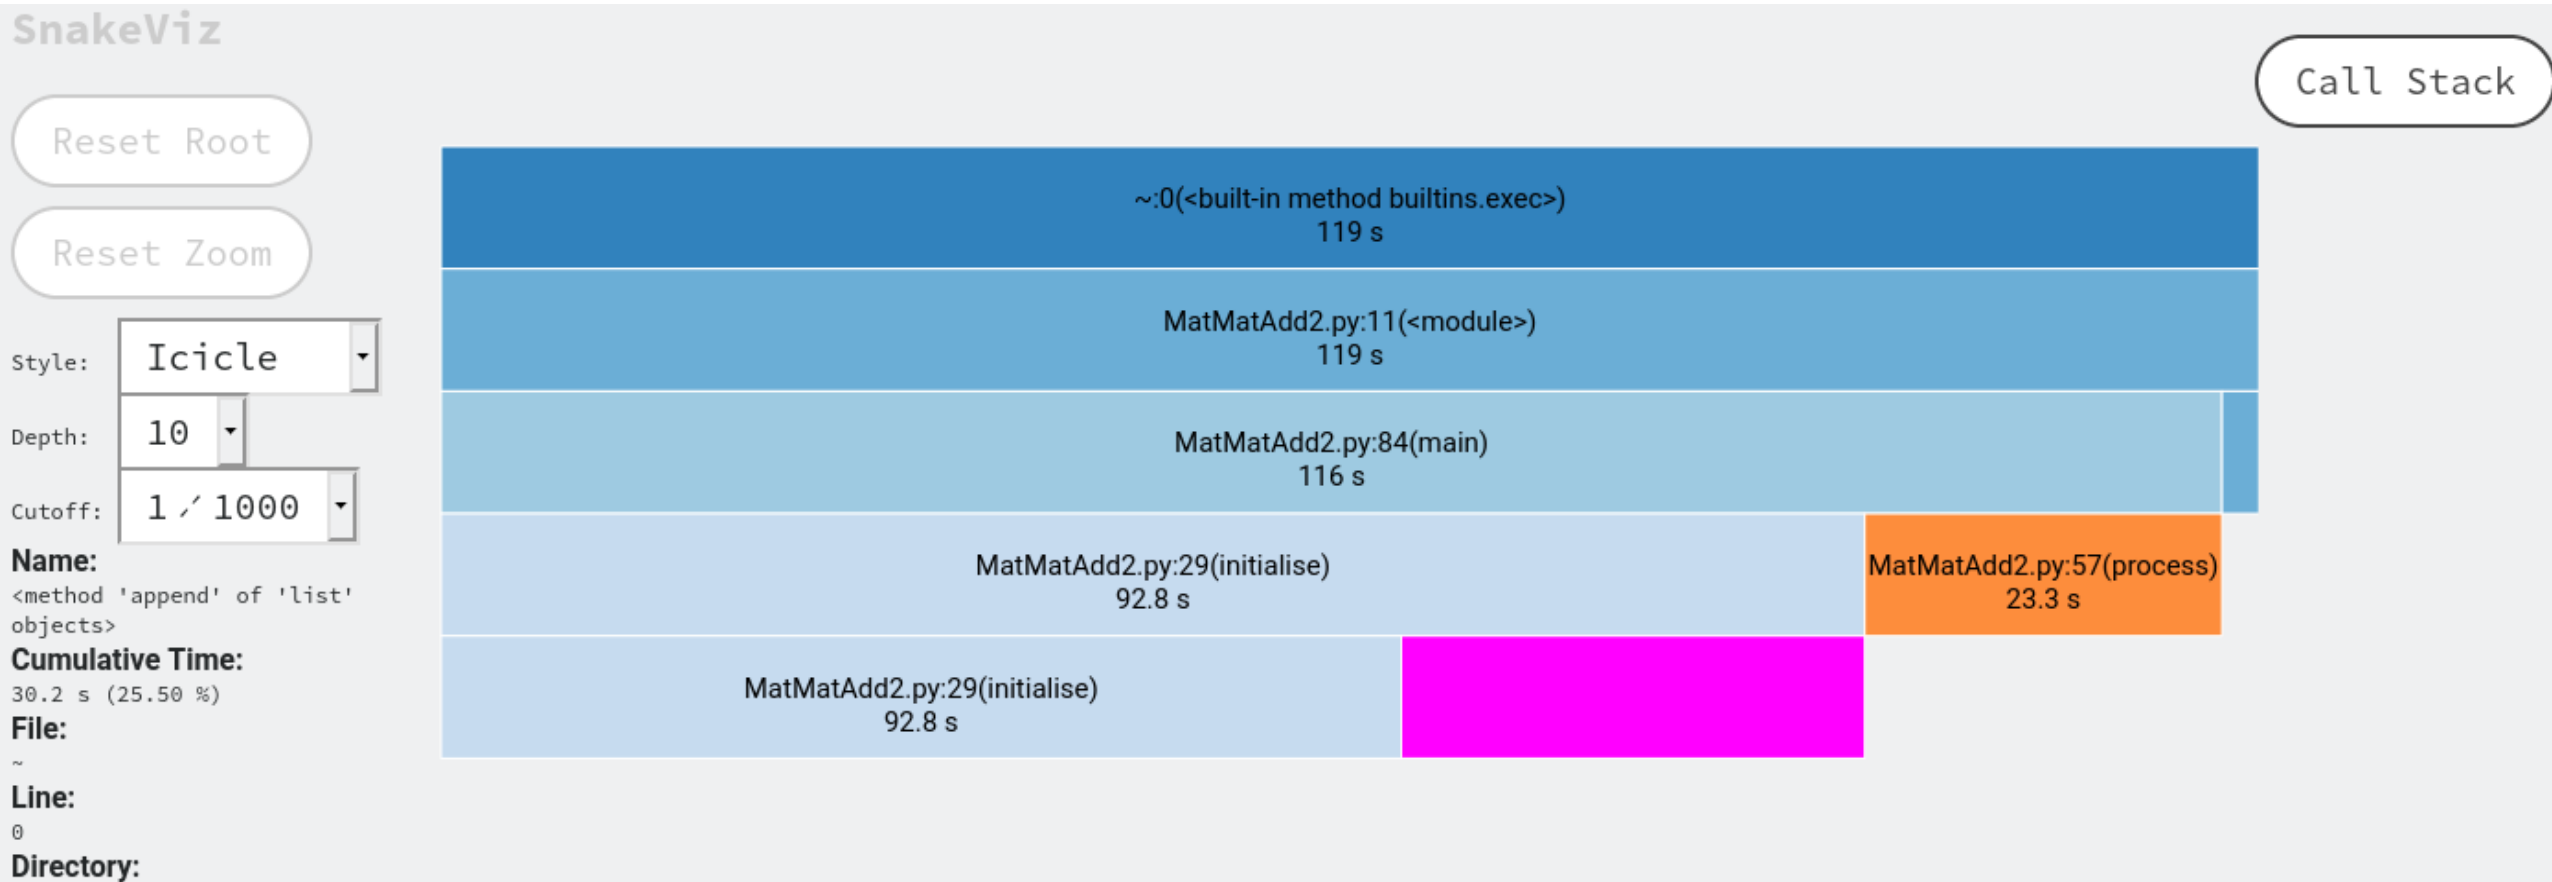
---

805355529 function calls in 118.819 seconds

Ordered by: cumulative time

	ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
	1	0.000	0.000	118.819	118.819	{built-in method builtins.exec}
	1	2.420	2.420	118.819	118.819	MatMatAdd2.py:11(<module>)
	1	0.000	0.000	116.398	116.398	MatMatAdd2.py:84(main)
	1	62.740	62.740	93.062	93.062	MatMatAdd2.py:29(initialise)
805355520		30.322	0.000	30.322	0.000	{method 'append' of 'list' objects}
	1	23.337	23.337	23.337	23.337	MatMatAdd2.py:57(process)
	1	0.000	0.000	0.000	0.000	MatMatAdd2.py:15(setup)
	1	0.000	0.000	0.000	0.000	{built-in method builtins.print}
	1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
	1	0.000	0.000	0.000	0.000	{built-in method builtins.len}

# Profiling Code (SnakeViz sample)



# Profiling Code (SnakeViz sample)

Search: <input type="text"/>					
ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	62.61	62.61	92.83	92.83	MatMatAdd2.py:29(initialise)
805355520	30.22	3.753e-08	30.22	3.753e-08	~:0(<method 'append' of 'list' objects>)
1	23.29	23.29	23.29	23.29	MatMatAdd2.py:57(process)
1	2.418	2.418	118.5	118.5	MatMatAdd2.py:11(<module>)
1	9.224e-06	9.224e-06	9.224e-06	9.224e-06	~:0(<built-in method builtins.print>)
1	5.393e-06	5.393e-06	116.1	116.1	MatMatAdd2.py:84(main)
1	3.409e-06	3.409e-06	1.294e-05	1.294e-05	MatMatAdd2.py:15(setup)
1	2.357e-06	2.357e-06	118.5	118.5	~:0(<built-in method builtins.exec>)
1	5.72e-07	5.72e-07	5.72e-07	5.72e-07	~:0(<method 'disable' of '_Isprof.Profiler' objects>)
1	3.07e-07	3.07e-07	3.07e-07	3.07e-07	~:0(<built-in method builtins.len>)

Showing 1 to 10 of 10 entries

# Line profiling

---

- Using the `line_profiler` package:

```
@profile          1. Add an annotation  
                  (@profile decorator)  
                  to your functions  
def MyFunction () :
```

```
kernprof -l -v MyPythonCode.py
```

2. Instrument your code  
for line profiling

```
python -m line_profiler MyPythonCode.py.lprof
```

3. Pass the instrumented code in the `MyPythonCode.py.lprof` script file  
to the `line_profiler` package

# Profiling Code (line\_profiler sample)

Timer unit: 1e-06 s

Total time: 494.288 s

File: MatMatAdd3.py

Function: initialise at line 29

Line #	Hits	Time	Per Hit	% Time	Line Contents
29					=====( Initialisation )=====
30					
31					@profile
32					def initialise ( matA, matB, matY, Rows, Cols ):
33					
34	16385	4557.0	0.3	0.0	# Initialize matrix A
35	16384	1079428.0	65.9	0.2	
36	268451840	75533955.0	0.3	15.3	for r in range ( Rows ): # A for loop for row entries
37	268435456	83015741.0	0.3	16.8	v = []
38	16384	5588.0	0.3	0.0	for c in range ( Cols ): # A for loop for column entries
39					v.append ( 1.0 )
40					matA.append ( v )
41	16385	4713.0	0.3	0.0	
42	16384	1270204.0	77.5	0.3	# Initialize matrix B
43	268451840	78478761.0	0.3	15.9	for r in range ( Rows ): # A for loop for row entries
44	268435456	86529825.0	0.3	17.5	v = []
45	16384	5662.0	0.3	0.0	for c in range ( Cols ): # A for loop for column entries
46					v.append ( 2.0 )
47					matB.append ( v )
48	16385	5157.0	0.3	0.0	
49	16384	1270371.0	77.5	0.3	# Initialize matrix Y
50	268451840	79801429.0	0.3	16.1	for r in range ( Rows ): # A for loop for row entries
51	268435456	87277163.0	0.3	17.7	v = []
52	16384	5913.0	0.4	0.0	for c in range ( Cols ): # A for loop for column entries
53					v.append ( 0.0 )
54	1	1.0	1.0	0.0	matY.append ( v )
55					
56					return ( matA, matB, matY )

# Memory consumption

---

- Using the `memory_profiler` package:

```
@profile          1. Add an annotation  
                  (@profile decorator)  
                  to your functions  
def MyFunction () :
```

```
python -m memory_profiler MyPythonCode.py
```

2. Pass the `MyPythonCode.py` script file  
to the `memory_profiler` package



# Profiling Code (memory\_profiler sample)

---

Filename: MatMatAdd5.py

Line #	Mem usage	Increment	Occurences	Line Contents
=====				
28	70.195 MiB	70.195 MiB	1	@profile
29				def initialise ( matA, matB, matY, Rows, Cols ):
30				
31				# Matrix shape is NxN ( N, N )
32	2118.391 MiB	2048.195 MiB	1	matA = np.full ( ( Rows, Cols ), 1.0, dtype = float )
33	4166.324 MiB	2047.934 MiB	1	matB = np.full ( ( Rows, Cols ), 2.0, dtype = float )
34				
35	4166.324 MiB	0.000 MiB	1	return ( matA, matB, matY )

Filename: MatMatAdd5.py

Line #	Mem usage	Increment	Occurences	Line Contents
=====				
39	4166.324 MiB	4166.324 MiB	1	@profile
40				def process ( matA, matB, matY ):
41				
42	6214.254 MiB	2047.930 MiB	1	matY = matA + matB
43	6214.254 MiB	0.000 MiB	1	return ( matY )

# Pre-requisites

---

- `pip install snakeviz`
- `pip install line_profiler`
- `pip install memory_profiler`