

# Parallel Algorithm Design

Departamento de Electrónica y Ciencias de la Computación  
Facultad de Ingeniería y Ciencias  
Pontificia Universidad Javeriana – Cali

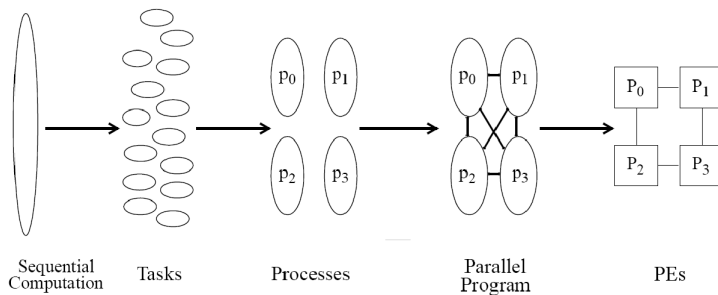
Santiago de Cali, Mayo de 2022



# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology
  - Partitioning
  - Communication
  - Agglomeration
  - Mapping
- 4 Case Studies
  - Boundary Value Problem
  - Finding the Maximum
  - The n-body Problem
- 5 Adding Data Input
- 6 WCET Calculation
- 7 Further Reading

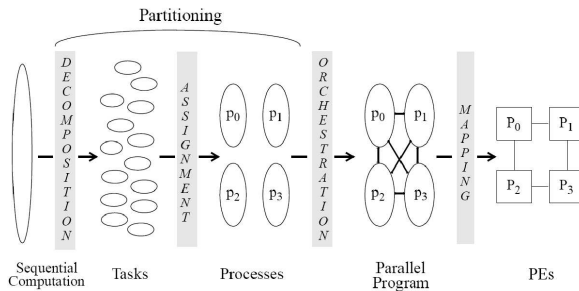
# Some Definitions<sup>1</sup>



- A **(Primitive) Task** is an arbitrarily defined piece of the work done by the program. It is the smallest unit of concurrency that the parallel program can exploit.
- A **Process** is an abstract entity that performs tasks. A **Parallel Program** is composed of multiple cooperating processes, each of which performs a subset of the tasks in the program. Processes may need to communicate and synchronise with one another to perform their assigned tasks.
- The way processes perform their assigned tasks is by executing them on the physical **PEs** in the machine.

<sup>1</sup>[Culler+98]


# Steps in the Process<sup>2</sup>



- 1 **Decomposition** of the Sequential Computation into Tasks.
- 2 **Assignment** of Tasks to Processes.
- 3 **Orchestration** of the necessary data accesses, communication, and synchronisation among Processes.
- 4 **Mapping** or binding of Processes to PEs.


<sup>2</sup>[Culler+98]

# Parallel Computation



Parallel Computation = set of *Processes*

# Process




## *Process*

- Program
- Local Memory
- Collection of I/O ports

# Parallel Computation

Can be represented by a directed graph:



- Vertices represent Processes
- Directed Edges represent *Communication Channels*
- Processes interact by sending messages through Communication Channels

# Communication Channel

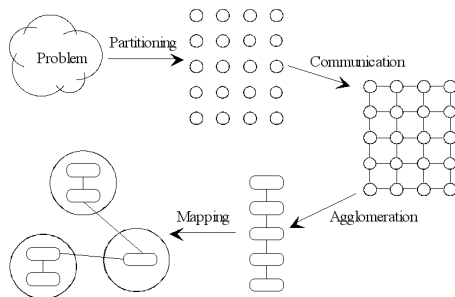
- Message Queue that connects one process output port with another process input port
- Receiving is a *synchronous* (blocking) operation
- Sending is an *asynchronous* (non-blocking) operation



# Execution Time

- The period of time during which **any process is active**
- The **starting time** is when the first process starts executing
- The **finishing time** is when the last process has stopped executing

# Foster's Parallel Algorithm Design Methodology<sup>3</sup> - Overview



- Together, **Decomposition** and **Assignment** are called **Partitioning**, since they divide the work done by the program among the cooperating processes.
- **Orchestration** is decomposed into **Communication** and **Agglomeration**.

---

<sup>3</sup>[Quinn04]

# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology**
  - **Partitioning**
  - Communication
  - Agglomeration
  - Mapping
- 4 Case Studies
- 5 Adding Data Input

# Foster's Parallel Algorithm Design Methodology (1)

## Since:

The number of primitive tasks is an upper bound on the maximum degree of parallelism one can exploit.


## Goal:

Identify as many primitive tasks as possible.


### 1 Partitioning (dividing computation and data into pieces):

- Domain decomposition (data-centric)
  - Divide data into pieces; e.g., an array into sub-arrays (reduction); a loop into sub-loops (matrix multiplication), a search space into sub-spaces (chess)
  - Determine how to associate computations with the data
- Functional decomposition (computation-centric)
  - Divide computation into pieces; e.g., pipelines (floating point multiplication), workflows (payroll processing)
  - Determine how to associate data with the computations

# Domain Decomposition



# Domain Decomposition Example



- Typically the focus is to divide the largest and/or the most frequently accessed data structure in the program
- It is usually best to maximise the number of primitive tasks

# Domain Decomposition Types

1D



BLOCK



CYCLIC

2D




BLOCK, \*



\*, BLOCK




BLOCK, BLOCK



CYCLIC, \*



\*, CYCLIC



CYCLIC, CYCLIC

## Domain Decomposition Types (2)

- Block
  - $n$  is multiple of  $p$ : straightforward. . .
  - $n$  is not multiple of  $p$ : Every process gets  $\lfloor \frac{n}{p} \rfloor$ ; process 0 gets the remainder  $(n \bmod p)$
- Cyclic (a.k.a. interleaved)
  - Process 0 is responsible for blocks  $0, 0 + p, 0 + 2p, \dots$
  - Process 1 is responsible for blocks  $1, 1 + p, 1 + 2p, \dots$
  - and so on. . .



## Domain Decomposition Types (3)

Block ( $n$  is not multiple of  $p$ ): Calculate  $r = n \bmod p > 0$

① Concentrate all of the larger blocks among the smaller-numbered processes:

- First  $r$  processes get a block of size  $\lceil \frac{n}{p} \rceil$
- The remaining  $p - r$  processes get a block of size  $\lfloor \frac{n}{p} \rfloor$
- First element controlled by process  $p\_id$ :  $p\_id \lfloor \frac{n}{p} \rfloor + \min(p\_id, r)$
- Last element controlled by process  $p\_id$ :  $(p\_id + 1) \lfloor \frac{n}{p} \rfloor + \min(p\_id + 1, r) - 1$

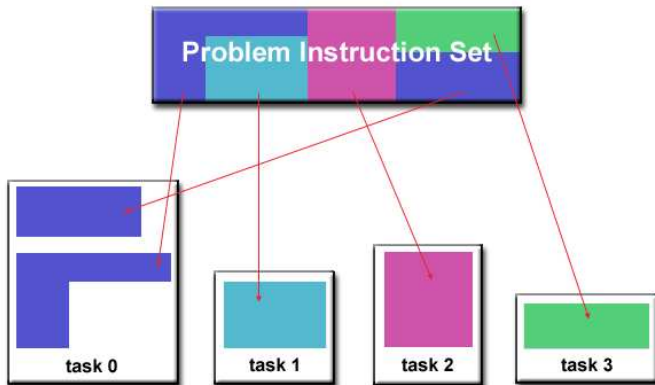
## Domain Decomposition Types (4)

Block ( $n$  is not multiple of  $p$ ): Calculate  $r = n \bmod p > 0$


② Distribute the larger blocks among the processes:

- First element controlled by process  $p\_id$ :  $\left\lfloor p\_id \frac{n}{p} \right\rfloor$
- Last element controlled by process  $p\_id$ :  $\left\lfloor (p\_id + 1) \frac{n}{p} \right\rfloor - 1$

# Functional Decomposition




# Functional Decomposition Example



- Task1: Track instrument positions for the next image
- Task2: Convert an image from physical coordinates to image coordinates
- Task3: Display the previous image

## Functional Decomposition Example 2



- The atmosphere model generates wind velocity data that are used by the ocean model
- The ocean model generates sea surface temperature data that are used by the atmosphere model
- and so on

## Foster's Parallel Algorithm Design Methodology (2)

### Since:

The number of primitive tasks is an upper bound on the maximum degree of parallelism one can exploit.

### Goal:

Identify as many primitive tasks as possible.

### Step 1:

Partitioning (dividing computation and data into pieces) **Checklist:**

- Are redundant computations/data storage minimised?
- Are primitive tasks roughly the same size?
- Is the number of tasks an increasing function of the problem size?

# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology**
  - Partitioning
  - Communication**
  - Agglomeration
  - Mapping
- 4 Case Studies
- 5 Adding Data Input

# Foster's Parallel Algorithm Design Methodology (3)

Since:

Communication between tasks is part of the overhead of a parallel algorithm.

Goal:

Minimise parallel overhead.

## 2 Communication (establish communication pattern):

- Local communication
  - Task needs values from a small number of other tasks
  - Create channels illustrating data flow
- Global communication
  - Significant number of tasks contribute data to perform a computation
  - Do not create channels for them early in design



## Foster's Parallel Algorithm Design Methodology (4)

### Since:

Communication between tasks is part of the overhead of a parallel algorithm.

### Goal:

Minimise parallel overhead.

### Step 2:

Communication (establish communication pattern)

### Checklist:

- Are communication operations balanced among tasks?
- Does each task communicates with only small group of neighbours?
- Can tasks perform communications concurrently?
- Can tasks perform computations concurrently?

# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology**
  - Partitioning
  - Communication
  - Agglomeration**
  - Mapping
- 4 Case Studies
- 5 Adding Data Input


# Foster's Parallel Algorithm Design Methodology (5)

- ③ Agglomeration (reducing the amount of parallel overhead):
  - Grouping tasks into larger tasks
  - **Goals**
    - Lower communication overhead (increase locality)
    - Maintain scalability (for porting purposes)
    - Reduce software engineering costs (by using existing sequential code)
  - In MPI programming, goal often to create one agglomerated task (process) per PE

# Agglomeration Examples



Combining tasks into processes eliminates communication and increases locality



Combining groups of sending and receiving tasks into processes reduces the number of messages being sent (fewer, longer messages)

# Foster's Parallel Algorithm Design Methodology (6)

## Step 3:

Agglomeration (reducing the amount of parallel overhead)

## Checklist:

- Does the locality increased?
- Do replicated computations take less time than the communication they replaced?
- Is the amount of replicated data small enough to allow the algorithm to scale?
- Do agglomerated tasks have similar computational and communications costs?

# Foster's Parallel Algorithm Design Methodology (7)

## Step 3:

Agglomeration (reducing the amount of parallel overhead)

## Checklist (cont.):

- Is the number of tasks an increasing function of the problem size?
- Is the number of tasks as small as possible, yet at least as great as the number of PEs in the likely target computers?
- Is the trade-off between the chosen agglomeration and the cost of modifications to existing sequential code reasonable?

# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology**
  - Partitioning
  - Communication
  - Agglomeration
  - Mapping**
- 4 Case Studies
- 5 Adding Data Input

# Foster's Parallel Algorithm Design Methodology (8)

## 4 Mapping (assigning tasks/processes to PEs):


- **Goals**

- Maximise PE utilisation
- Minimise interPE communication
- Centralised multiprocessor: mapping done by operating system
- Distributed memory system: mapping done by user

Finding an optimal solution is NP-hard. . .



# Mapping Example



- What if every PE has the same speed and that every task/process requires the same amount of execution time?
- What if every channel communicates the same amount of data?

# Mapping Decision Tree

- Static number of processes
  - Structured communication
    - Constant computation time per task  
Agglomerate tasks into processes to minimise communication  
Create one process per PE
    - Variable computation time per task  
Cyclically map tasks/processes to PEs
  - Unstructured communication
    - Use a static load-balancing algorithm
- Dynamic number of tasks/processes

# Mapping Decision Tree

- Static number of processes
- Dynamic number of tasks/processes
  - Use a run-time task-scheduling algorithm
    - centralised:  
e.g., a master-slave strategy
    - distributed:  
push strategy: PEs with too many processes send some of them to neighbouring PEs  
pull strategy: PEs with no work to do ask neighbouring PEs for work
  - Use a dynamic load-balancing algorithm  
e.g., share load among neighbouring PEs; remapping periodically

# Foster's Parallel Algorithm Design Methodology (9)

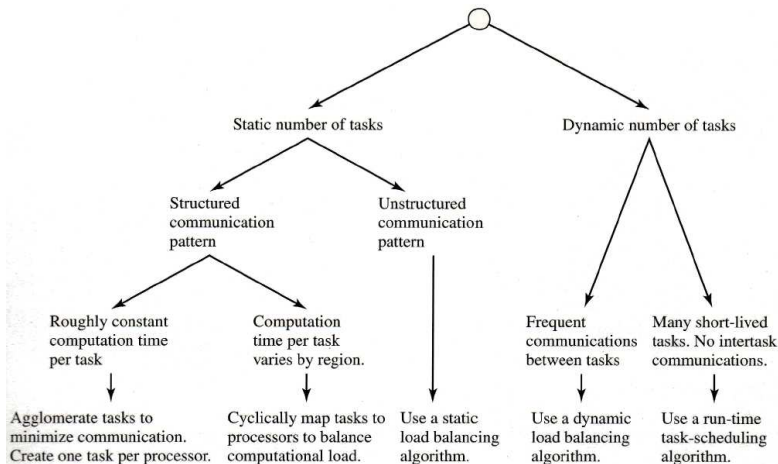
## Step 4:

Mapping (assigning processes to PEs)

## Checklist:

- Considered designs based on one process per PE and multiple processes per PE
  - If multiple processes per PE chosen, ratio of processes to PEs is at least 10:1
- Evaluated static and dynamic process allocation
- If dynamic process allocation chosen, process allocator is not a bottleneck to performance


# Mapping Decision Tree



# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology
- 4 Case Studies**
  - **Boundary Value Problem**
  - Finding the Maximum
  - The n-body Problem
- 5 Adding Data Input


# Introduction



- Rod length = 1
- Ice bath temperature = 0 °C
- $temp_0(x) = 100 \sin(\pi x)$

# Rod Cools as Time Progresses

- A partial differential equation models the temperature at any point of the rod at any point in time




A finite difference approximation to the rod-cooling problem

- Each curve represents the temperature distribution of the rod at some point in time




# Finite Difference Approximation – Data Structure (1)



Every point  $u_{i,j}$  represents a matrix element containing the temperature at position  $i$  on the rod at time  $j$

- The rod is divided into  $n$  sections of length  $h$ , so each row has  $n + 1$  elements. Increasing  $n$  reduces the error in the approximation
- Time from 0 to  $T$  is divided into  $m$  discrete entities of length  $k$ , so the matrix contains  $m + 1$  rows

# Finite Difference Approximation – Data Structure (2)



- The finite difference algorithm steps forward in time, using values from time  $j$  to compute the value for time  $j + 1$  using the formula

$$u_{i,j+1} = ru_{i-1,j} + (1 - 2r)u_{i,j} + ru_{i+1,j}$$


where  $r = k/h^2$

# Partitioning

- One data item per grid point
- Associate one primitive task with each grid point
- Two-dimensional domain decomposition


# Communication

Identify communication pattern between primitive tasks:



- Each primitive task computing  $u_{i,j+1}$  requires the values of  $u_{i-1,j}$ ,  $u_{i,j}$ , and  $u_{i+1,j}$
- Each interior primitive task has three incoming and three outgoing channels
- Primitive tasks on the edges have fewer channels

# Agglomeration



Tasks computing rod temperatures later in time depend upon the results produced by tasks computing rod temperatures earlier in time

- Agglomerate all the tasks associated with each point in the rod into a process



A linear array of processes

# Mapping



A linear array of processes

In a real problem the number of rod segments would be large

- Agglomerate processes so that computational workloads are balanced and communication is minimised



A process is responsible for computing, over all time steps, the temperatures for a contiguous group of rod locations

# Sequential Execution Time

- $t_{comp} = \tau$  – time to compute  $u_{i,j+1}$
- $n$  – number of pieces of size  $h$
- $m$  – number of time steps
- $t_s = m(n - 1)t_{comp}$

# Parallel Execution Time

- $p$  – number of PEs
- $\gamma$  – time to communicate a value to another task
- $t_p = t_{comp} + t_{comm}$ 
  - $t_{comp} = \left\lceil \frac{n-1}{p} \right\rceil \tau$  – computation time for each iteration
  - $t_{comm} = 2\gamma + 2\gamma$  – time to send/receive to/from left and right neighbours for each iteration
- $t_p = m \times \left( \left\lceil \frac{n-1}{p} \right\rceil \tau + 4\gamma \right)$



# Outline

- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology
- 4 Case Studies**
  - Boundary Value Problem
  - Finding the Maximum**
  - The n-body Problem
- 5 Adding Data Input

# Introduction (1)

Boundary value problems arising from the real world are too complicated to solve analitically

- $x$  – computed solution
- $c$  – correct solution
- $error = \left| \frac{x - c}{c} \right|$

Computed	0.15	0.16	0.16	0.19
Correct	0.15	0.16	0.17	0.18
Error (%)	0.00%	0.00%	6.25%	5.26%

Maximum error = 6.25 %

# Introduction (2)

Given

- a set of values  $a_0, a_1, a_2, \dots, a_{n-1}$
- associative operator  $\oplus$

**reduction** is the process of computing

$$a_0 \oplus a_1 \oplus a_2 \dots \oplus a_{n-1}$$

Examples:

- Add
- Multiply
- And, Or
- Maximum, Minimum

# Partitioning

Since reduction requires exactly  $n - 1$  operations, it has  $\Theta(n)$  time complexity on a sequential computer. How quickly can we perform a reduction on a parallel computer?

- Divide the list into  $n$  pieces
- Associate one primitive task with each piece

# Communication (1)

Brute force approach:




One task (the root task) receives a list element from each of the other  $n - 1$  tasks and performs all the additions

- $\gamma$  – time to communicate a value to another task
- $\tau$  – time to perform an addition
- $t_p = (n - 1)(\gamma + \tau)$  – slower than the sequential algorithm...

## Communication (2)


What if two tasks cooperate to perform the reduction?



- In time  $\left(\frac{n}{2} - 1\right)(\gamma + \tau)$  each semiroot task has a subtotal for its half of the elements
- In one additional communication/computation step a single task has the grand total
- $t_p = \frac{n}{2}(\gamma + \tau)$


## Communication (3)

Why not continue the process?



- Each task is responsible for  $n/4$  of the list elements
- After the four subtotals have been computed, two remaining communication/computation steps yield the grand total

# Communication – Binomial Trees




$$k = 0$$

- In a tree with  $p = 2^k$  nodes, the maximum distance from any node to the root in the lower left corner is  $k = \log_2 p$
- The **binomial tree** (a subgraph of a **hypercube**) is one of the most common communication patterns in parallel algorithm design




# Communication – Binomial Trees



$k = 1$

- In a tree with  $p = 2^k$  nodes, the maximum distance from any node to the root in the lower left corner is  $k = \log_2 p$
- The **binomial tree** (a subgraph of a **hypercube**) is one of the most common communication patterns in parallel algorithm design


# Communication – Binomial Trees



$k = 2$

- In a tree with  $p = 2^k$  nodes, the maximum distance from any node to the root in the lower left corner is  $k = \log_2 p$
- The **binomial tree** (a subgraph of a **hypercube**) is one of the most common communication patterns in parallel algorithm design


# Communication – Binomial Trees



$k = 3$

- In a tree with  $p = 2^k$  nodes, the maximum distance from any node to the root in the lower left corner is  $k = \log_2 p$
- The **binomial tree** (a subgraph of a **hypercube**) is one of the most common communication patterns in parallel algorithm design


# Communication – Binomial Trees



$$k = 3$$


- In a tree with  $p = 2^k$  nodes, the maximum distance from any node to the root in the lower left corner is  $k = \log_2 p$
- The **binomial tree** (a subgraph of a **hypercube**) is one of the most common communication patterns in parallel algorithm design

# Communication – Using Binomial Trees




16 tasks cooperate to find the summation value

# Communication – Using Binomial Trees




Half of the tasks send values, and half of the tasks receive values

# Communication – Using Binomial Trees



Half of the tasks add values, and half of the tasks become inactive


# Communication – Using Binomial Trees



A quarter of the tasks send values, and a quarter of the tasks receive values




# Communication – Using Binomial Trees




A quarter of the tasks add values, and three quarters of the tasks are inactive

# Communication – Using Binomial Trees




An eighth of the tasks send values, and an eighth of the tasks receive values

# Communication – Using Binomial Trees




An eighth of the tasks add values, and seven eighths of the tasks are inactive

# Communication – Using Binomial Trees




One task send its value, and one task receive a value

# Communication – Using Binomial Trees



One task add two values, and the remaining 15 tasks are inactive

# Communication – Using Binomial Trees




The task/channel graph forms a binomial tree

# Agglomeration and Mapping

The number of tasks is static, computations per task are trivial, and the communication pattern is regular (agglomerate tasks to minimise communication by following the previous graph)


- 16 tasks are mapped to 4 PEs



# Agglomeration and Mapping

The number of tasks is static, computations per task are trivial, and the communication pattern is regular (agglomerate tasks to minimise communication by following the previous graph)

- 4 tasks on each PE are agglomerated into a single process






# Agglomeration and Mapping

The number of tasks is static, computations per task are trivial, and the communication pattern is regular (agglomerate tasks to minimise communication by following the previous graph)


- Each PE adds  $n/p$  values



# Agglomeration and Mapping

The number of tasks is static, computations per task are trivial, and the communication pattern is regular (agglomerate tasks to minimise communication by following the previous graph)


- Two processes send values, and two processes receive values and add



# Agglomeration and Mapping

The number of tasks is static, computations per task are trivial, and the communication pattern is regular (agglomerate tasks to minimise communication by following the previous graph)


- One process send its value, and one process receives a value and add. The latter has the grand total



# Outline


- 1 Parallelisation Process
- 2 The Process/Channel Programming Model
- 3 Foster's Design Methodology
- 4 Case Studies**
  - Boundary Value Problem
  - Finding the Maximum
  - **The n-body Problem**
- 5 Adding Data Input

# Introduction (1)




- Every particle exerts a gravitational pull on every other particle
- The white particle has a particular position and velocity vector (indicated by the white arrow)
- Its future position is influenced by the gravitational forces exerted by the other two particles


# Introduction (2)



# Introduction (2)




# Introduction (2)





# Introduction (3)




Solved by performing computations on all pairs of bodies

- gravitational interaction =  $G \frac{m_1 m_2}{d^2}$
- typically, time complexity =  $\Theta(n^2)$ , where  $n$  is the number of bodies

## Introduction (4)

A group of bodies that is far enough away from a given body may be approximated by the centre of mass of the group. The farther apart the bodies, the larger the group that may be thus approximated:




# Partitioning

- Domain partitioning
- Assume one task per particle
- Task has particle's position, velocity vector
- Iteration
  - Get positions of all other particles
  - Compute new position, velocity


# Communication (1)

A **gather** operation is a global communication that takes a dataset distributed among a group of processes and collects the items on a single process:




# Communication (1)

A **gather** operation is a global communication that takes a dataset distributed among a group of processes and collects the items on a single process:




## Communication (2)

An **all-gather** operation is similar to gather, except at the end of the communication every process has a copy of the entire dataset:




## Communication (2)

An **all-gather** operation is similar to gather, except at the end of the communication every process has a copy of the entire dataset:



## Communication (3)

Since it is necessary to update the location of every particle, an **all-gather** operation is called for:




Set up a channel between every pair of tasks:

- $p - 1$  communication steps
- message length: 2 coordinates



## Communication (3)

Since it is necessary to update the location of every particle, an **all-gather** operation is called for:



Use a bottom-up approach:

- $\log_2 p$  communication steps
- message length:  $2 \times 2^{i-1}$ , where  $i$  is the exchange step

# Agglomeration and Mapping

In general, there are far more particles  $n$  than PEs  $p$ .

- Assume  $n$  is a multiple of  $p$ .
- Associate one process per PE and agglomerate  $n/p$  particles into each process.
- The all-gather communication operation requires  $\log_2 p$  communication steps.
- In the first step the messages have length  $n/p$ , in the second step the messages have length  $2n/p$ , etc.

# Communication Time

- Setting up a channel between every pair of processes:


$$\begin{aligned}t_{comm} &= (p - 1) \times \left( \lambda + \frac{n/p}{\beta} \right) \\&= \lambda(p - 1) + \frac{n(p - 1)}{\beta p}\end{aligned}$$

- Using all-gather (for each iteration):

$$\begin{aligned}t_{comm} &= \sum_{i=1}^{\log_2 p} \left( \lambda + \frac{2^{i-1} n}{\beta p} \right) \\&= \lambda \log_2 p + \frac{n(p - 1)}{\beta p}\end{aligned}$$

# Introduction

- Assume a single process (0) is responsible for performing file I/O operations
- By adding new channels for file I/O the resulting process/channel graph is:




- A pair of coordinates identify a particle's location
- Velocity can be represented by another pair of values
- Reading the position and velocities of  $n$  particles requires time

$$t_{read} = \lambda_{IO} + \frac{4n}{\beta_{IO}}$$


## Communication (1)

After the I/O process inputs the particles, it is necessary to break up the input data into pieces so that each process has its assigned subsection containing  $n/p$  elements. This global operation is called **scatter**:




## Communication (1)

After the I/O process inputs the particles, it is necessary to break up the input data into pieces so that each process has its assigned subsection containing  $n/p$  elements. This global operation is called **scatter**:




## Communication (2)

One way to scatter the particles is for the I/O process to simply send the corresponding  $n/p$  particles to each of the other tasks. Another is to derive a scatter operation requiring  $\log_2 p$  steps:



## Communication (2)


One way to scatter the particles is for the I/O process to simply send the corresponding  $n/p$  particles to each of the other tasks. Another is to derive a scatter operation requiring  $\log_2 p$  steps:





## Communication (2)

One way to scatter the particles is for the I/O process to simply send the corresponding  $n/p$  particles to each of the other tasks. Another is to derive a scatter operation requiring  $\log_2 p$  steps:




## Scattering Time

- The I/O process sends the corresponding  $n/p$  particles to each of the other tasks:


$$\begin{aligned}t_{\text{scattering}} &= (p - 1) \times \left( \lambda + \frac{4n}{\beta p} \right) \\&= \lambda(p - 1) + \frac{4n}{\beta p} \times (p - 1)\end{aligned}$$

- Using a scatter operation requiring  $\log_2 p$  steps:

$$\begin{aligned}t_{\text{scattering}} &= \sum_{i=1}^{\log_2 p} \left( \lambda + \frac{4n}{2^i \beta p} \right) \\&= \lambda \log_2 p + \frac{4n}{\beta p} \times \left( 1 - \frac{1}{p} \right)\end{aligned}$$




$$\sum_{i=0}^{s-1} WCET(V_i)$$



$$\max \{ WCET(path_j) \}$$

$$\forall j = 1, \dots, a$$



$$I \times \sum_{i=0}^{s-1} WCET(V_i)$$

-  Cormen T. H., Leiserson C. E., Rivest R. L., Stein C. Introduction to Algorithms. Second Edition. MIT Press, 2002 (Appendix A – Summations)
-  Culler D., Singh J. P., Gupta, A. Parallel Computer Architecture – A Hardware-Software Approach. Morgan Kaufmann, 1998
-  Quinn M. J. Parallel Programming in C with MPI and OpenMP. McGraw Hill, 2004